

Location-Based Species Presence Prediction

Michael Ward
tue97607@temple.edu

April 28, 2022

Contents

1	Abstract	1
2	Introduction	2
2.1	Background	2
2.2	Data Description and Visualization	3
2.2.1	Observations	3
2.2.2	Environmental Vectors	5
2.2.3	Satellite Imagery	7
2.3	Accuracy Evaluation	7
3	Related Work	8
3.1	Mapping Species Distributions	8
3.2	Review of Machine Learning Based SDM	9
3.3	Predicting Rare Species Distributions	10
4	Methodology	11
5	Discussion	11
5.1	Results	11
5.2	Limitations	11
6	Acknowledgments	11
7	References	11

1 Abstract

This paper

2 Introduction

2.1 Background

Distribution and diversity of plant and animal species have both always varied in space and time. However, the rise in human population and activities in the last few centuries have drastically impacted the speed and magnitude of change in species distribution and diversity [10]. Understanding where different species live is crucial to making well-informed conservation decisions. The rapid change in distribution and diversity of species has made the process of conservation decision making harder, since species location data quickly loses its value over time. While humans have caused major disruptions in the distribution and diversity of species, we also generate millions of geo-coded species observations every year through the work of citizen science projects, which can be used to build species distribution models (commonly referred to as SDM) [8].

The aim of this project is to address the need for predictive species distributions with models based on geo-localized observations of plant and animal species from France and the United States. The task is to return a set of 30 candidate species that should contain the true observed species when provided with a GPS position and its associated features. This work is relevant to the aforementioned fields of biodiversity management and conservation. Specifically, this type of SDM could allow for biodiversity inventories via the development of location-based recommendation services, and accelerate the recording and validation of species observations for the purposes of creating of large-scale and high-quality data sets. Ultimately, the desire is to be able to predict what types of species might be found at a new location, which can be used for the purpose of conservation decision-making.

This project has an associated Kaggle competition. The competition is held jointly as part of the LifeCLEF 2022 lab and FGVC9 workshop. While this has been an ongoing yearly competition since 2018, there is still considerable room for improvement for this problem. The evaluation metric for this problem is the top-30 error rate, which will be discussed further in section 2.3. I am submitting my work on this project to the Kaggle competition, so I will be evaluating my models according to the top-30 error rate metric that will be used in the competition to evaluate model performance.

The winning solution in 2021 for GeoLifeClef still had a very high top-30 error rate, near 75%, and was not able to make use of all of the covariate data available. According to an overview of the 2021 competition results, the main challenges associated with this problem have to do with properly aggregating the different covariates,

and determining how much complementary or redundant information is contained in the high resolution patches compared to the ones captured from the bioclimatic and soil variables [8]. I intend to survey current and past research on species prediction to inform my testing of multiple different models that attempt to solve this problem.

2.2 Data Description and Visualization


2.2.1 Observations

The dataset contains roughly 1.6 million observations of plant and animal species. There are about 17,000 species in the dataset, 9,000 of which are plant species, while the other 8,000 are animal species. Each observation contains the species name and the gps coordinates where it was observed. Each observation is also paired with a set of covariates which contain characteristics of the surrounding landscape and environment. The covariates are in tabular, .TIFF, .JPEG, and .GeoTIFF format, and are a combination of high resolution remote sensing imagery, land cover data, altitude, as well as low-resolution climate and soil variables. The observations in this dataset are heavily imbalanced. As seen in Figure 1, the minimum number of observations per class is 3, while the max number of observations per class is 6,701. Additionally shown below in Figure 2 is a graph illustrating the imbalance on a logarithmic scale to emphasize how many classes can be considered as minority classes in this dataset. Over 30% of the 17,037 classes have just 10 observations or less, out of the roughly 1.6 million total observations.

Figure 1: Dataset Metrics

Number of Classes (Unique Species)	17,037
Number of Observations (Training)	1,627,475
Number of Observations (Validation)	40,080
Number of Observations (Testing)	36,421
Max Number of Observations per Class	6,701
Min Number of Observations per Class	3

Figure 2: observation distribution



The observations in this data set have been recorded both from the United States and from France. Shown below are the spatial distributions of observations in each country.

Figure 3: US spatial distribution

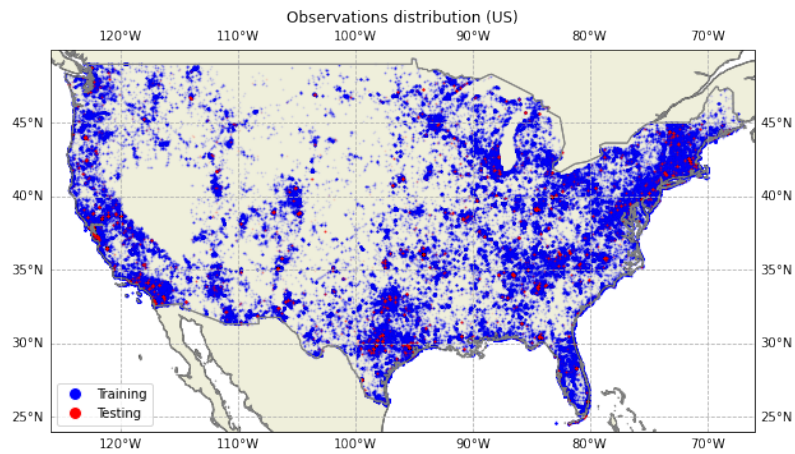



Figure 4: FR spatial distribution



The train/test split for the data has already been determined by the competition, and uses a spatial block holdout procedure to split the data. The reasoning for using a spatial block holdout is to limit the effects of spatial auto-correlation on the data. The way this splitting procedure works is by assigning observations to a grid of 5km by 5km quadrants, of which 2.5% are randomly sampled to create the test set, another 2.5% randomly sampled to create the validation set, and the remaining are used for the training set. The validation set is created using the same spatial block holdout procedure [8].

2.2.2 Environmental Vectors

The main group of data provided by GeoLifeClef is made up of environmental vectors (of which there are 27). These environmental vectors were sourced from two separate databases, WorldClim and SoilGrids. WorldClim is a database containing high spatial resolution global weather and climate data [6]. 19 of the 27 environmental vectors provided by GeoLifeClef are sourced from WorldClim, and mostly includes various Bio-Climatic variables related to temperature and precipitation. SoilGrids is another database of global digital soil mapping that uses machine learning methods to map the spatial distribution of soil properties around the world [5]. The 19

bio-climatic variables sourced from WorldClim were extracted from satellite raster imagery with a spatial resolution of $1\text{km}^2/\text{pixel}$. The 8 pedologic (soil) variables were extracted from satellite raster imagery with a spatial resolution of $250\text{m}^2/\text{pixel}$ [8].

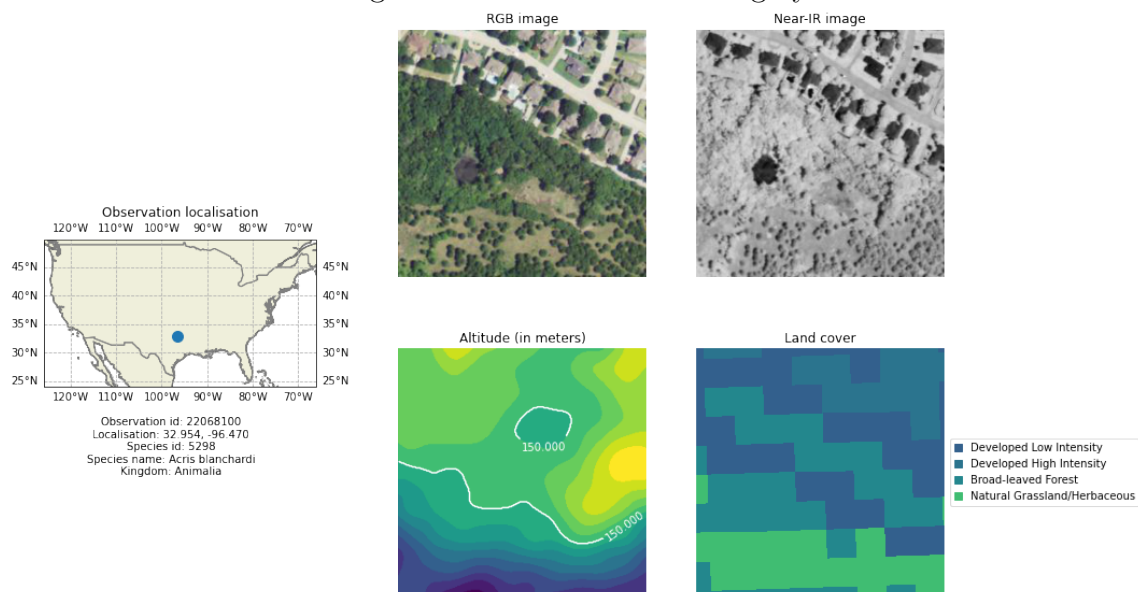
Figure 5: Environmental Vectors

Bio-Climatic Variables (WorldClim) [6]	
Name	Description
bio_1	Annual Mean Temperature
bio_2	Mean Diurnal Range (Mean of monthly (max temp - min temp))
bio_3	Isothermality (bio_2/bio_7) (* 100)
bio_4	Temperature Seasonality (standard deviation *100)
bio_5	Max Temperature of Warmest Month
bio_6	Min Temperature of Coldest Month
bio_7	Temperature Annual Range (bio_5-bio_6)
bio_8	Mean Temperature of Wettest Quarter
bio_9	Mean Temperature of Driest Quarter
bio_10	Mean Temperature of Warmest Quarter
bio_11	Mean Temperature of Coldest Quarter
bio_12	Annual Precipitation
bio_13	Precipitation of Wettest Month
bio_14	Precipitation of Driest Month
bio_15	Precipitation Seasonality (Coefficient of Variation)
bio_16	Precipitation of Wettest Quarter
bio_17	Precipitation of Driest Quarter
bio_18	Precipitation of Warmest Quarter
bio_19	Precipitation of Coldest Quarter
Pedologic Variables (SoilGrids) [5]	
Name	Description
orcdrc	Soil organic carbon content (g/kg at 15cm depth)
phihox	Ph x 10 in H2O (at 15cm depth)
cecsol	cation exchange capacity of soil in cmolc/kg 15cm depth
bdticm	Absolute depth to bedrock in cm
clyppt	Clay (0-2 micro meter) mass fraction at 15cm depth
sltppt	Silt mass fraction at 15cm depth
sndppt	Sand mass fraction at 15cm depth
bldfie	Bulk density in kg/m3 at 15cm depth

2.2.3 Satellite Imagery

Each observation in the dataset has 4 associated high resolution satellite images, one covering the rgb spectrum, one covering the near infrared spectrum (NIR), one with altitude measurements, and one with land cover classifications. Each image are a spatial resolution of 1 meter per pixel, and are 256 x 256 pixels (256 meters x 256 meters). Each image is centered on it's associated observation. The United States RGB and NIR images were sourced from the 2009 to 2011 cycle of the National Agriculture Imagery Program (NAIP), while the French RGB and NIR images are sourced from the BD-ORTH 2.0 and ORTHO-HR 1.0 databases from the Institut national de l'information géographique et forestière. The United States land cover imagery that is used in this dataset is sourced from the National Land Cover Database, and the French land cover imagery is sourced from the Centre d'Etudes Spatiales de la Biosphère [7]. The elevation imagery is sourced from the NASA Shuttle Radar Topography Mission for both the US and French observations. Below is an example of the satellite raster imagery available for an observation.

Figure 6: Satellite Raster Imagery



2.3 Accuracy Evaluation

The main metric that I will use to evaluate my models is the top-30 error rate, which is a set-valued classification metric defined by the Kaggle competition. Each

observation i in the dataset is associated with a single ground-truth label y_i which corresponds to the observed species. The model must provide 30 candidate labels $\hat{y}_{i,1}, \hat{y}_{i,2}, \dots, \hat{y}_{i,30}$ for each observation. The top 30 error rate is then computed using the following equation:

$$\text{Top-30 error rate} = \frac{1}{N} \sum_{i=1}^N e_i \quad \text{where} \quad e_i = \begin{cases} 1 & \text{if } \forall k \in \{1, \dots, 30\}, \hat{y}_{i,k} \neq y_i \\ 0 & \text{otherwise} \end{cases}$$

There are several reasons to use this type of error measure as the main evaluation metric for this problem and data. The first reason being that this data set contains presence-only observations, meaning that each observation only contains the presence of one species whilst many different species could potentially be recorded from each geographic location. A set-valued classification metric allows for more precision when measuring the accuracy of predicted sets when dealing with the limited knowledge inherent with presence-only observations. Second, recent research has shown that using single-output predictions for ambiguous multi-class datasets leads to overall poor performance and results [1]. For a dataset such as this with over 17,000 classes and many classes with as few as 3 observations, it makes sense to use a more lenient set-valued classification metric which allows for more meaningful comparison of models and input data types, instead of a single-output prediction per observation.

3 Related Work

3.1 Mapping Species Distributions

The most widely cited resource when it comes to the general topic of species distribution modeling is a 2017 textbook named Mapping Species Distributions by Janet Franklin. This text is a comprehensive overview on the topic of species distribution modeling, covering not just the models used within the field, but also the history and ecological basis of species distribution modeling. This text breaks the SDM methods into two broad categories of statistical methods, and machine learning based methods. Among the statistical methods discussed in this text are generalized linear models (GLM), generalized additive models (GAM), and multivariate adaptive regression splines (MARS).

This text identifies generalized linear modeling as the most established / most generally accepted statistical framework for SDM. Their reasoning for the popularity

of GLM in SDM is the ability of the link function to handle response variables with irregular distributions. Overall, the author is critical of the statistical approaches to SDM due to their assumption that observations are independent. These approaches violate Tobler’s ‘First Law of Geography’, that near things tend to be similar to each other due to the fact that they are likely to influence each other or be influenced by the same local environmental factors [4]. I am in agreement with the author that this law of geography should not be ignored when working with spatial data.

Franklin also covers the most commonly used machine learning techniques used for SDM at the time, which are identified broadly as ensemble methods, deep learning methods, and genetic algorithms. Out of these categories, the methods that have been most widely successful in the field of ecology and for the purposes of SDM are random forests and other ensemble methods such as stochastic gradient boosting, various flavors of neural networks, and support vector machines. Ensemble methods are identified as the most applied method, due to their great success with complex classification problems. While neural networks and SVM have also shown promise in this field, the author shows evidence that while these methods have performed well on complex classification problems, they sometimes performed worse or no better than statistical methods like GLMs or GANs, or ensemble methods like random forests. In addition to this, neural networks have a steep learning curve to be used effectively which has resulted in a slower adoption in the field of ecology compared to other machine learning based methods [4].

3.2 Review of Machine Learning Based SDM

Another 2017 publication talks in more detail about the specific benefits and flaws of several machine learning techniques with respect to their usage in species distribution modeling. These authors surveyed 258 journal and conference papers published between 2000 and 2017 that contained both ‘machine learning’ and ‘species’ distribution’ to determine what the most common machine learning methods in SDM are today. In this paper, the authors compare random forests, MaxEnt, support vector machines, and artificial neural networks. These four methods are not exhaustive, but comprise the four most widely used methods out of the surveyed material.

The authors provide an overview of cases where these methods were applied, and give their thoughts on the usage of each method. First, the authors highlight the importance of combining these methods with expert domain knowledge: their research found that the usefulness of these machine learning methods are highly dependent on the ecological assumptions that influenced the choice of environmental variables used. Ultimately, the ideal relationship should be one where ecologists focus

on ecological explainability, while data scientists focus on predictive performance.

In terms of model evaluation in SDM, the authors group random forests, MaxEnt, and SVM together as ensemble methods and compare them to the usage of deep learning. The authors find that ensemble methods are the most widely applied methods in SDM, and have achieved considerable results. In comparison, deep learning methods have been slower to adopt in the ecology field, though the authors have found successful application of deep learning models in tree species classification, plant disease detection, and plant phenotyping. They cite three reasons for deep learning having less prevalence in the field of ecology, comparatively to ensemble methods; First, they agree with Franklin that deep learning is more difficult and time consuming for ecologists to learn and master. Additionally, it is less preferred due to its lack of ecological interpretability. Lastly, the hardware or cloud platforms necessary for supporting deep learning when working with large amounts of raster data is quite expensive. Despite these reasons for deep learning’s lack of current usage in SDM and ecology, the authors believe that adoption will increase over time due to the increase of software platforms that support deep learning, and the possibility of deep neural networks performing better than ensemble methods on larger and more complex data sets [12].

3.3 Predicting Rare Species Distributions

One of the underlying issues of SDMs has to do with generalizing species distributions in large undersampled areas, especially with few observations available [9]. This is certainly a problem for this project, as over 30% of the ~17,000 species have 10 observations or less, and the spatial distribution of observations is not consistent across either the United States or France. One paper has assessed what methods can best address this underlying issue of SDMs by testing multiple machine learning models on three case studies containing rare species distributions in China [9]. Other research, such as the papers and textbooks previously mentioned, has proved the effectiveness of machine learning techniques in ecology and SDM, but has not tested which methods perform best with imbalanced observations and rare classes [2].

Chunrong Mi et al. quantify which models out of stochastic gradient boosting, random forest, CART, and MaxEnt, perform best on three separate cases of rare species distribution in China (Hooded Crane with 33 observations, White Naped Crane with 40 observations, and Black Necked Crane with 75 observations) [9]. This study was most relevant to my methodology out of the other literature found due to both the addressing of rare species distribution, and the usage of WorldClim environmental vectors for the predictive variables. Chunrong Mi et al. used WorldClim

variables bio_1 through bio_15 in addition to several variables obtained from [naturalearthdata.com](https://natural-earth.org/) which provides distance to relevant land and water features for their predictive variables. This was relevant, as I am using WorldClim variables bio_1 through bio_19 in my models.

4 Methodology

methods here

5 Discussion

5.1 Results

5.2 Limitations

6 Acknowledgments

I would like to thank the contributors and maintainers of the GeoLifeCLEF Github repository: Maxi Miliense, Titouan Lorieul, Benjamin Deneu, Chris Botella, and Elijah Cole. This repository contains open source Python code that was designed for GeoLifeCLEF Kaggle challenge competitors, specifically for loading, plotting, and submitting the data provided in the competition. This code was especially useful for me to get up and running quickly with loading and visualizing the data, so I could move on to transforming the data and running models. I used the provided python functions for spatially plotting data, and generating the competition submission files. I would also like to thank Temple’s CIS 4523/5523 Professor Zoran Obradovic and teaching assistant Marija Stanojevic for providing valuable feedback on this project during project presentations.

7 References

- [1] Evgenii Chzhen et al. “Set-valued classification—overview via a unified framework”. In: *arXiv preprint arXiv:2102.12318* (2021).
- [2] D Richard Cutler et al. “Random forests for classification in ecology”. In: *Ecology* 88.11 (2007), pp. 2783–2792.

- [3] Benjamin Deneu et al. “Location-based species recommendation using co-occurrences and environment-GeoLifeCLEF 2018 challenge”. In: *Working Notes of CLEF 2018-Conference and Labs of the Evaluation Forum*. Vol. 2125. 2018.
- [4] Janet Franklin. *Mapping species distributions: spatial inference and prediction*. Cambridge University Press, 2010.
- [5] Tomislav Hengl et al. “SoilGrids250m: Global gridded soil information based on machine learning”. In: *PLoS one* 12.2 (2017), e0169748.
- [6] Robert J Hijmans et al. “Very high resolution interpolated climate surfaces for global land areas”. In: *International Journal of Climatology: A Journal of the Royal Meteorological Society* 25.15 (2005), pp. 1965–1978.
- [7] Collin Homer et al. “Completion of the 2011 National Land Cover Database for the conterminous United States—representing a decade of land cover change information”. In: *Photogrammetric Engineering & Remote Sensing* 81.5 (2015), pp. 345–354.
- [8] Titouan Lorieul et al. “Overview of GeoLifeCLEF 2021: Predicting species distribution from 2 million remote sensing images”. In: *Working Notes of CLEF 2021-Conference and Labs of the Evaluation Forum*. Vol. 2936. 2021, pp. 1451–1462.
- [9] Chunrong Mi et al. “Why choose Random Forest to predict rare species distribution with few samples in large undersampled areas? Three Asian crane species models provide supporting evidence”. In: *PeerJ* 5 (2017), e2849.
- [10] Jesús N Pinto-Ledezma and Jeannine Cavender-Bares. “Predicting species distributions and community composition using satellite remote sensing predictors”. In: *Scientific Reports* 11.1 (2021), pp. 1–12.
- [11] Sachith Seneviratne. “Contrastive representation learning for natural world imagery: Habitat prediction for 30,000 species”. In: *CLEF working notes* (2021).
- [12] Jian Zhang and Sen Li. “A Review of Machine Learning Based Species’ Distribution Modelling”. In: *2017 International Conference on Industrial Informatics - Computing Technology, Intelligent Technology, Industrial Information Integration (ICIICII)*. 2017, pp. 199–206. DOI: 10.1109/ICIICII.2017.76.