

Location-Based Species Presence Prediction

Michael Ward
tue97607@temple.edu

April 28, 2022

Contents

1	Abstract	1
2	Introduction	2
2.1	Background	2
2.2	Data Description and Visualization	3
2.2.1	Observations	3
2.2.2	Environmental Vectors	4
2.2.3	Satellite Imagery	5
2.3	Accuracy Evaluation	6
3	Related Work	6
3.1	Mapping Species Distributions	6
3.2	Review of Machine Learning Based SDM	7
3.3	Predicting Rare Species Distributions	8
4	Methodology	9
4.1	Missing Data	9
4.2	Baseline	10
4.3	Modeling	10
4.3.1	KNN	10
4.3.2	Random Forests	11
4.3.3	Deep Neural Network	11
5	Discussion	12
5.1	Results	12
5.2	Limitations	13
6	Acknowledgments	14
7	References	15

1 Abstract

This paper

2 Introduction

2.1 Background

Distribution and diversity of plant and animal species have both always varied in space and time. However, the rise in human population and activities in the last few centuries have drastically impacted the speed and magnitude of change in species distribution and diversity [10]. Understanding where different species live is crucial to making well-informed conservation decisions. The rapid change in distribution and diversity of species has made the process of conservation decision making harder, since species location data quickly loses its value over time. While humans have caused major disruptions in the distribution and diversity of species, we also generate millions of geo-referenced species observations every year through the work of citizen science projects, which can be used to build species distribution models (commonly referred to as SDM) [8].

The aim of this project is to address the need for predictive species distributions with models based on geo-localized observations of plant and animal species from France and the United States. The task is to return a set of 30 candidate species that should contain the true observed species when provided with a GPS position and its associated features. This work is relevant to the aforementioned fields of biodiversity management and conservation. Specifically, this type of SDM could allow for biodiversity inventories via the development of location-based recommendation services, and accelerate the recording and validation of species observations for the purposes of creating of large-scale and high-quality data sets. Ultimately, the desire is to be able to predict what types of species might be found at a new location, which can be used for the purpose of conservation decision-making.

This project has an associated Kaggle competition. The competition is held jointly as part of the LifeCLEF 2022 lab and FGVC9 workshop. While this has been an ongoing yearly competition since 2018, there is still considerable room for improvement for this problem. The evaluation metric for this problem is the top-30 error rate, which will be discussed further in section 2.3. I am submitting my work on this project to the Kaggle competition, so I will be evaluating my models according to the top-30 error rate metric that will be used in the competition to evaluate model performance.

The winning solution in 2021 for GeoLifeClef still had a very high top-30 error rate, near 75%, and was not able to make use of all of the covariate data available. According to an overview of the 2021 competition results, the main challenges associated with this problem have to do with properly aggregating the different covariates,


and determining how much complementary or redundant information is contained in the high resolution patches compared to the ones captured from the bioclimatic and soil variables [8]. I intend to survey current and past research on species prediction to inform my testing of multiple different models that attempt to solve this problem.

2.2 Data Description and Visualization

2.2.1 Observations

The dataset contains roughly 1.6 million observations of plant and animal species. There are about 17,000 species in the dataset, 9,000 of which are plant species, while the other 8,000 are animal species. Each observation contains the species name and the gps coordinates where it was observed. Each observation is also paired with a set of covariates which contain characteristics of the surrounding landscape and environment. The covariates are in tabular, .TIFF, .JPEG, and .GeoTIFF format, and are a combination of high resolution remote sensing imagery, land cover data, altitude, as well as low-resolution climate and soil variables. The observations in this dataset are heavily imbalanced. The minimum number of observations per class is 3, while the max number of observations per class is 6,701. Shown below in Figure 1 is a graph illustrating the imbalance on a logarithmic scale to emphasize how many classes can be considered as minority classes in this dataset. Over 30% of the 17,037 classes have just 10 observations or less, out of the roughly 1.6 million total observations.

Figure 1: observation distribution



The observations in this data set have been recorded both from the United States and from France. Shown below are the spatial distributions of observations in each

country.

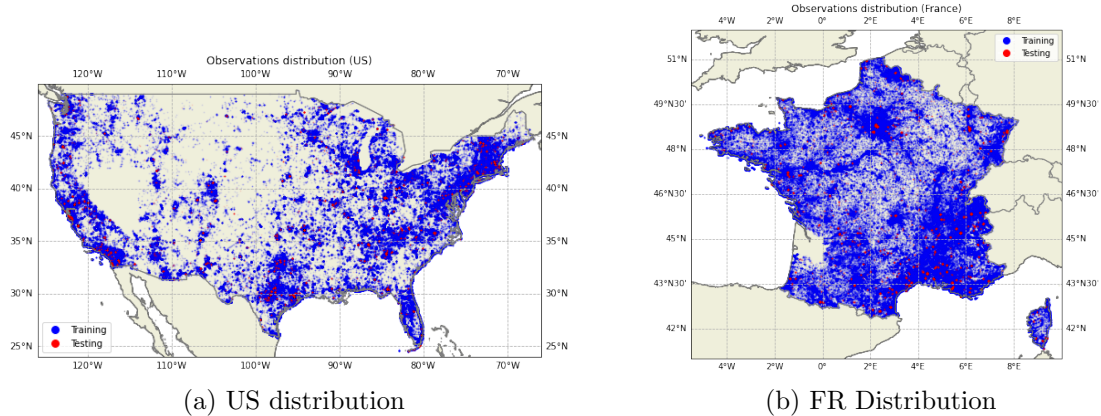


Figure 2: Spatial Distribution of Observations

The train/test split for the data has already been determined by the competition, and uses a spatial block holdout procedure to split the data. The reasoning for using a spatial block holdout is to limit the effects of spatial auto-correlation on the data. The way this splitting procedure works is by assigning observations to a grid of 5km by 5km quadrants, of which 2.5% are randomly sampled to create the test set, another 2.5% randomly sampled to create the validation set, and the remaining are used for the training set. The validation set is created using the same spatial block holdout procedure [8].

2.2.2 Environmental Vectors

The main group of data provided by GeoLifeClef is made up of environmental vectors (of which there are 27). These environmental vectors were sourced from two separate databases, WorldClim and SoilGrids. WorldClim is a database containing high spatial resolution global weather and climate data [6]. 19 of the 27 environmental vectors provided by GeoLifeClef are sourced from WorldClim, and mostly includes various Bio-Climatic variables related to temperature and precipitation. SoilGrids is another database of global digital soil mapping that uses machine learning methods to map the spatial distribution of soil properties around the world [5]. The 19 bio-climatic variables sourced from WorldClim were extracted from satellite raster imagery with a spatial resolution of $1\text{km}^2/\text{pixel}$. The 8 pedologic (soil) variables were extracted from satellite raster imagery with a spatial resolution of $250\text{m}^2/\text{pixel}$ [8].

Figure 3: Environmental Vectors

Bio-Climatic Variables (WorldClim) [6]	
Name	Description
bio_1	Annual Mean Temperature
bio_2	Mean Diurnal Range (Mean of monthly (max temp - min temp))
bio_3	Isothermality (bio_2/bio_7) (* 100)
bio_4	Temperature Seasonality (standard deviation *100)
bio_5	Max Temperature of Warmest Month
bio_6	Min Temperature of Coldest Month
bio_7	Temperature Annual Range (bio_5-bio_6)
bio_8	Mean Temperature of Wettest Quarter
bio_9	Mean Temperature of Driest Quarter
bio_10	Mean Temperature of Warmest Quarter
bio_11	Mean Temperature of Coldest Quarter
bio_12	Annual Precipitation
bio_13	Precipitation of Wettest Month
bio_14	Precipitation of Driest Month
bio_15	Precipitation Seasonality (Coefficient of Variation)
bio_16	Precipitation of Wettest Quarter
bio_17	Precipitation of Driest Quarter
bio_18	Precipitation of Warmest Quarter
bio_19	Precipitation of Coldest Quarter
Pedologic Variables (SoilGrids) [5]	
Name	Description
orcdrc	Soil organic carbon content (g/kg at 15cm depth)
phihox	Ph x 10 in H2O (at 15cm depth)
cecsol	cation exchange capacity of soil in cmolc/kg 15cm depth
bdticm	Absolute depth to bedrock in cm
clyppt	Clay (0-2 micro meter) mass fraction at 15cm depth
sltppt	Silt mass fraction at 15cm depth
sndppt	Sand mass fraction at 15cm depth
bldfie	Bulk density in kg/m3 at 15cm depth

2.2.3 Satellite Imagery

Each observation in the dataset has 4 associated high resolution satellite images, one covering the rgb spectrum, one covering the near infrared spectrum (NIR), one with altitude measurements, and one with land cover classifications. Each image are a spatial resolution of 1 meter per pixel, and are 256 x 256 pixels (256 meters x 256 meters). Each image is centered on it's associated observation. The United States RGB and NIR images were sourced from the 2009 to 2011 cycle of the National Agriculture Imagery Program (NAIP), while the French RGB and NIR images are sourced from the BD-ORTH 2.0 and ORTHO-HR 1.0 databases from the Institut national de l'information géographique et forestière. The United States land cover imagery that is used in this dataset is sourced from the National Land Cover Database, and the French land cover imagery is sourced from the Centre d'Etudes Spatiales de la Biosphère [7]. The elevation imagery is sourced from the NASA Shuttle Radar Topography Mission for both the US and French observations.

2.3 Accuracy Evaluation

The main metric that I will use to evaluate my models is the top-30 error rate, which is a set-valued classification metric defined by the Kaggle competition. Each observation i in the dataset is associated with a single ground-truth label y_i which corresponds to the observed species. The model must provide 30 candidate labels $\hat{y}_{i,1}, \hat{y}_{i,2}, \dots, \hat{y}_{i,30}$ for each observation. The top 30 error rate is then computed using the following equation:

$$\text{Top-30 error rate} = \frac{1}{N} \sum_{i=1}^N e_i \quad \text{where} \quad e_i = \begin{cases} 1 & \text{if } \forall k \in \{1, \dots, 30\}, \hat{y}_{i,k} \neq y_i \\ 0 & \text{otherwise} \end{cases}$$

There are several reasons to use this type of error measure as the main evaluation metric for this problem and data. The first reason being that this data set contains presence-only observations, meaning that each observation only contains the presence of one species whilst many different species could potentially be recorded from each geographic location. A set-valued classification metric allows for more precision when measuring the accuracy of predicted sets when dealing with the limited knowledge inherent with presence-only observations. Second, recent research has shown that using single-output predictions for ambiguous multi-class datasets leads to overall poor performance and results [1]. For a dataset such as this with over 17,000 classes and many classes with as few as 3 observations, it makes sense to use a more lenient set-valued classification metric which allows for more meaningful comparison of models and input data types, instead of a single-output prediction per observation.

3 Related Work

3.1 Mapping Species Distributions

The most widely cited resource when it comes to the general topic of species distribution modeling is a 2017 textbook named Mapping Species Distributions by Janet Franklin. This text is a comprehensive overview on the topic of species distribution modeling, covering not just the models used within the field, but also the history and ecological basis of species distribution modeling. This text breaks the SDM methods into two broad categories of statistical methods, and machine learning based methods. Among the statistical methods discussed in this text are generalized lin-

ear models (GLM), generalized additive models (GAM), and multivariate adaptive regression splines (MARS).

This text identifies generalized linear modeling as the most established / most generally accepted statistical framework for SDM. Their reasoning for the popularity of GLM in SDM is the ability of the link function to handle response variables with irregular distributions. Overall, the author is critical of the statistical approaches to SDM due to their assumption that observations are independent. These approaches violate Tobler’s ‘First Law of Geography’, that near things tend to be similar to each other due to the fact that they are likely to influence each other or be influenced by the same local environmental factors [4]. I am in agreement with the author that this law of geography should not be ignored when working with spatial data.

Franklin also covers the most commonly used machine learning techniques used for SDM at the time, which are identified broadly as ensemble methods, deep learning methods, and genetic algorithms. Out of these categories, the methods that have been most widely successful in the field of ecology and for the purposes of SDM are random forests and other ensemble methods such as stochastic gradient boosting, various flavors of neural networks, and support vector machines. Ensemble methods are identified as the most applied method, due to their great success with complex classification problems. While neural networks and SVM have also shown promise in this field, the author shows evidence that while these methods have performed well on complex classification problems, they sometimes performed worse or no better than statistical methods like GLMs or GAMs, or ensemble methods like random forests. In addition to this, neural networks have a steep learning curve to be used effectively which has resulted in a slower adoption in the field of ecology compared to other machine learning based methods [4].

3.2 Review of Machine Learning Based SDM

Another 2017 publication talks in more detail about the specific benefits and flaws of several machine learning techniques with respect to their usage in species distribution modeling. These authors surveyed 258 journal and conference papers published between 2000 and 2017 that contained both ‘machine learning’ and ‘species’ distribution’ to determine what the most common machine learning methods in SDM are today. In this paper, the authors compare random forests, MaxEnt, support vector machines, and artificial neural networks. These four methods are not exhaustive, but comprise the four most widely used methods out of the surveyed material.

The authors provide an overview of cases where these methods were applied, and give their thoughts on the usage of each method. First, the authors highlight

the importance of combining these methods with expert domain knowledge: their research found that the usefulness of these machine learning methods are highly dependent on the ecological assumptions that influenced the choice of environmental variables used. Ultimately, the ideal relationship should be one where ecologists focus on ecological explainability, while data scientists focus on predictive performance.

In terms of model evaluation in SDM, the authors group random forests, Max-Ent, and SVM together as ensemble methods and compare them to the usage of deep learning. The authors find that ensemble methods are the most widely applied methods in SDM, and have achieved considerable results. In comparison, deep learning methods have been slower to adopt in the ecology field, though the authors have found successful application of deep learning models in tree species classification, plant disease detection, and plant phenotyping. They cite three reasons for deep learning having less prevalence in the field of ecology, comparatively to ensemble methods; First, they agree with Franklin that deep learning is more difficult and time consuming for ecologists to learn and master. Additionally, it is less preferred due to its lack of ecological interpretability. Lastly, the hardware or cloud platforms necessary for supporting deep learning when working with large amounts of raster data is quite expensive. Despite these reasons for deep learning’s lack of current usage in SDM and ecology, the authors believe that adoption will increase over time due to the increase of software platforms that support deep learning, and the possibility of deep neural networks performing better than ensemble methods on larger and more complex data sets [12].

3.3 Predicting Rare Species Distributions

One of the underlying issues of SDMs has to do with generalizing species distributions in large undersampled areas, especially with few observations available [9]. This is certainly a problem for this project, as over 30% of the 17,000 species have 10 observations or less, and the spatial distribution of observations is not consistent across either the United States or France. One paper has assessed what methods can best address this underlying issue of SDMs by testing multiple machine learning models on three case studies containing rare species distributions in China [9]. Other research, such as the papers and textbooks previously mentioned, has proved the effectiveness of machine learning techniques in ecology and SDM, but has not tested which methods perform best with imbalanced observations and rare classes [2].

Chunrong Mi et al. quantify which models out of stochastic gradient boosting, random forest, CART, and MaxEnt, perform best on three separate cases of rare species distribution in China (Hooded Crane with 33 observations, White Naped

Crane with 40 observations, and Black Necked Crane with 75 observations) [9]. This study was most relevant to my methodology out of the other literature found due to both the addressing of rare species distribution, and the usage of WorldClim environmental vectors for the predictive variables. Chunrong Mi et al. used WorldClim variables bio_1 through bio_15 in addition to several variables obtained from naturalearthdata.com which provides distance to relevant land and water features for their predictive variables. This was relevant, as I am using WorldClim variables bio_1 through bio_19 in my models.

The authors found that random forests had the best performance in predicting the distribution of each of the three rare crane species by a significant amount, using area under ROC curve and the true skill statistic as accuracy metrics. The authors also comment that the application of a few observations to project a distribution area beyond its sampled range is a large challenge that hasn't been attempted until recently in ecology and SDM. Because of this, they recommend that modelers should assess their model performance on not just an evaluation metric (such as AUC or TSS), but also combine the use of visualization and expert domain knowledge to assess how the predicted species distribution map looks and compares to real data [9].


4 Methodology

4.1 Missing Data

Roughly 3.6% of observations are missing data for the soil variables sourced from SoilGrids, and roughly 2.6% of observations are missing data for the geo-climatic variables sourced from WorldClim. Before making any decisions on how to handle the missing data I plotted a missingness graph to see how much data was missing, and whether there were correlations in the missing data.

The missingness plot below shows which variables are missing data on the x-axis, and shows the position in the data set in which they are missing on the y-axis. If a line runs straight across the x-axis, that means that there is a correlation in the missing data, and it should not be considered missing at random. In this case, you can see that all of the missing data is either correlated to WorldClim or SoilGrid - if an observation is missing one WorldClim variable it is missing all WorldClim variables. Likewise, if an observation is missing one SoilGrid variable, it is missing all other SoilGrid variables. A small portion of missing data (roughly 1%) is 100% correlated, meaning that it is missing all variables.

Figure 4: Missingness Plot



Since the missing data is missing not at random and many minority classes have less than 10 observations, it is not appropriate to drop the missing rows of data in this case. I decided to test my models with two different methods of imputation. The first method of imputation I used was a form of simple imputation, where the mean value of a variable is imputed for that variable's missing values. The second form of imputation I tested was k-nearest neighbors imputation, where each missing value is imputed using the mean value of that observation's k-nearest neighbors. I used a value of 20 for k for knn imputation.

4.2 Baseline

In order to have a baseline to compare my model results to, besides other competitors in the Kaggle competition, I went with a relatively simple baseline that makes sense for the dataset that I am working with. As a reminder, the top-30 error rate metric requires 30 species to be predicted for each observation in the dataset. I created a baseline that just estimates the top-30 most present species in the dataset for each observation. This baseline has a top-30 error rate of 93.5%.

4.3 Modeling

4.3.1 KNN

One machine learning classifier that I hadn't seen come up frequently during my review of related literature is the k-nearest neighbors classifier. I decided to give this classifier a try to see how it would compare specifically to random forests on this data.

While testing KNN, I iteratively increased K until it reached diminishing returns on the top 30 error rate. After reaching the rate of diminishing returns, I used grid search with 5 fold cross validation to tune the KNN hyper parameters. I found the optimal k-value to be 35 for this data. After finding an optimal k-value, I specifically tested different distance metrics, prediction weights, and the KNN algorithm used to compute nearest neighbors (KDTree, BallTree, and brute force search). I found from hyper parameter tuning that for this model, minkowski distance, uniform distance weighting, and the BallTree algorithm were the best fitting parameters. After tuning hyper parameters, I tested the model using several variations of the dataset. I tested it using simple imputation for missing features vs. knn imputation, and using normalized vs. non-normalized features. I ran each of the KNN models using all 27 environmental variables as features.

4.3.2 Random Forests

The second classification method I used in my methodology was random forests. All of the literature I reviewed pointed to random forests as the best technique for species distribution modeling, especially when working with rare and sparse species. For this method I also used grid search with 5 fold cross validation to tune the hyper parameters of the model and test multiple iterations. The specific parameters that I tested were differing numbers of trees per forest, the splitting criterion used (Gini vs. entropy), and the maximum tree depth allowed. I found from this testing that the best parameters for my model are random forests of 80 trees, a max tree depth of 16, and the use of Gini impurity instead of entropy for splitting. After tuning the hyper parameters of the model, I tested the model with several variations of the training / testing data. I tested versions of the data that had mean imputation vs. knn imputation applied, as well as with normalized vs. non-normalized features. It isn't necessary to normalize features for random forests, but since I had already prepared normalized data for use in a neural network I tested it here as well to see how it would effect the random forest model performance. I tested each random forest model using all 27 environmental variables as features.

4.3.3 Deep Neural Network

The last classification method that I wanted to use in my experiment was a deep neural network, though I ran into several roadblocks while attempting to test this method. My intention was to use a neural network with the provided raster imagery to predict the species observations. The first issue I encountered was when trying to convert the rasters to numpy arrays and feeding them into a multi-layer perceptron

classifier using Python’s scikit-learn module. Unfortunately, as I learned, scikit-learn is not able to take advantage of GPU acceleration for additional processing power. The amount of satellite imagery provided by the competition totals to over 1,000 gigabytes, and scikit-learn was not anywhere near capable of processing this much data with my available hardware. My second route was to learn a new framework for neural networks that could take advantage of hardware acceleration - I found that PyTorch would be a reasonable fit for this. However, due to the steep learning curve of PyTorch and time-restrictions, I was not able to get a model up and running in time to include in this report. The Kaggle competition will be open for another month, so I hope to be able to gain a better understanding of PyTorch in that time and finish composing a neural network to train on my data.

5 Discussion

5.1 Results

Figure 5: Model Results

Model Description	Top 30 error rate
KNN with mean imputation	90.3%
KNN with knn imputation	90.1%
KNN with knn impute and norm.	92.5%
KNN with mean imputatation and norm.	90.9%
Random Forest with mean imputation	77.4%
Random Forest with knn imputation	79.2%
Random Forest with knn impute and norm.	85.3%
Random Forest with mean imputatation and norm.	93.2%

Random forests had the best performance overall, when using data that was not normalized and missing values imputed with the mean. After reviewing the literature I had expected random forests to be the best performing model, since it has been widely applied for the purposes of species distribution modeling, and recently has been found to work well for modeling rare species distributions. Since I have one more month to submit additional results to the Kaggle competition, I plan to test some additional ensemble methods as well to see if better performance than my random forest model can be obtained with a different ensemble method. Based on my review of literature, I’d like to explore the application of support vector machines and stochastic gradient boosting on this problem.

I also want to talk about the effect of the different imputation methods on the results for both models. For KNN, knn imputation only lowered the error rate by 0.2% compared to mean imputation, while for Random Forests it actually increased the error rate by 1.8%. I was surprised to see that the more sophisticated imputation method was outperformed by mean imputation, but I have some thoughts as to why this occurred. Missing data in this data set is highly correlated among variables, so if an observation is missing data at all, it is either missing all bio-climatic variables, all soil variables, or a very small chance of missing all environmental variables. So when knn imputation is applied, an observation with missing data has very limited features to compare neighbors with - which I think is leading to poor performance of this imputer. This imputer is likely finding poor fitting neighbors for observations with missing features.

In my testing, KNN took significant tuning to surpass the simple baseline of the top-30 most present species. I think that the main reason for poor performance of this model with this specific dataset has to do with the fact that the data set is not only heavily imbalanced, but that a third of the classes have 10 observations or less. Because of this, I think it is likely that many of the minority classes have more observations from other classes nearby in the feature space than observations from their own class. So while KNN was able to slightly surpass the simple baseline, it doesn't seem that KNN is a good choice of model for species distribution modeling when the data set contains many rare classes.

5.2 Limitations

I want to begin by talking about a couple of limitations that I encountered which have to do with the qualities of the data set. The most obvious limitation of the data set is the high number of minority species: roughly 10% of the species have only 3 observations, and more than 30% of the species have 10 observations or less. The small number of observations for many species is a larger problem overall than the class imbalance, as this limits the potential accuracy of the model and also limits the use of methods that can be used to correct class imbalance such as SMOTE. I had wanted to do a run using SMOTE oversampling, but this was not appropriate given that 10% of the observations have a maximum of 2 neighbors to use for SMOTE oversampling. If more observations were able to be collected for the species that have the fewest observations, this would open up more possibilities for oversampling techniques, and likely greatly increase the accuracy of multiple models. Additionally, the data set is missing one feature that I think could be quite valuable for this type of research: a datetime feature. I was surprised to see that there was not a date time

feature to indicate the date and time that each observation was collected. For species distribution, this can hold valuable information because certain animal species are more likely to be observed at certain times of day and at certain periods of the year. Likewise, certain plant species are more or less likely to be observed at certain periods throughout the year. This type of feature would add noise to species that are observed regardless of time of day or time of year, but I think it would likely be a meaningful feature for identifying species that only appear at certain times of day or certain times of year.

Aside from limitations inherent in the data, I also encountered limitations based on my available hardware and skills. Due to the size of the tabular data (27 features x 1.6 million observations), scikit-learn was not able to run random forest models with parameters exceeding 100 trees per forest and depth greater than 20 without exceeding the available memory. Additionally, it was not able to run a multi-layer perceptron at all with either the tabular data, or the satellite imagery. Had I been more experienced in a framework that can take advantage of hardware acceleration at this point in time, like PyTorch, I may have been able to model the data using a neural network as well. My experience with attempting a neural network model for species distribution modeling lines up with the related work and literature that I reviewed, which found that the largest reason for lack of neural network usage in SDM is due to learning curve and hardware capabilities. Despite this, I would like to continue learning so that I can test this data on a neural network, as some other participants in the Kaggle competition this year have had great results so far in applying the provided satellite imagery to a CNN.

6 Acknowledgments

I would like to thank the contributors and maintainers of the GeoLifeCLEF Github repository: Maxi Miliense, Titouan Lorieul, Benjamin Deneu, Chris Botella, and Elijah Cole. This repository contains open source Python code that was designed for GeoLifeCLEF Kaggle challenge competitors, specifically for loading, plotting, and submitting the data provided in the competition. This code was especially useful for me to get up and running quickly with loading and visualizing the data, so I could move on to transforming the data and running models. I used the provided Python functions for spatially plotting data, and generating the competition submission files. I would also like to thank Temple's CIS 4523/5523 Professor Zoran Obradovic and teaching assistant Marija Stanojevic for providing valuable feedback on this project during project presentations.

7 References

- [1] Evgenii Chzhen et al. “Set-valued classification—overview via a unified framework”. In: *arXiv preprint arXiv:2102.12318* (2021).
- [2] D Richard Cutler et al. “Random forests for classification in ecology”. In: *Ecology* 88.11 (2007), pp. 2783–2792.
- [3] Benjamin Deneu et al. “Location-based species recommendation using co-occurrences and environment-GeoLifeCLEF 2018 challenge”. In: *Working Notes of CLEF 2018-Conference and Labs of the Evaluation Forum*. Vol. 2125. 2018.
- [4] Janet Franklin. *Mapping species distributions: spatial inference and prediction*. Cambridge University Press, 2010.
- [5] Tomislav Hengl et al. “SoilGrids250m: Global gridded soil information based on machine learning”. In: *PLoS one* 12.2 (2017), e0169748.
- [6] Robert J Hijmans et al. “Very high resolution interpolated climate surfaces for global land areas”. In: *International Journal of Climatology: A Journal of the Royal Meteorological Society* 25.15 (2005), pp. 1965–1978.
- [7] Collin Homer et al. “Completion of the 2011 National Land Cover Database for the conterminous United States—representing a decade of land cover change information”. In: *Photogrammetric Engineering & Remote Sensing* 81.5 (2015), pp. 345–354.
- [8] Titouan Lorieul et al. “Overview of GeoLifeCLEF 2021: Predicting species distribution from 2 million remote sensing images”. In: *Working Notes of CLEF 2021-Conference and Labs of the Evaluation Forum*. Vol. 2936. 2021, pp. 1451–1462.
- [9] Chunrong Mi et al. “Why choose Random Forest to predict rare species distribution with few samples in large undersampled areas? Three Asian crane species models provide supporting evidence”. In: *PeerJ* 5 (2017), e2849.
- [10] Jesús N Pinto-Ledezma and Jeannine Cavender-Bares. “Predicting species distributions and community composition using satellite remote sensing predictors”. In: *Scientific Reports* 11.1 (2021), pp. 1–12.
- [11] Sachith Seneviratne. “Contrastive representation learning for natural world imagery: Habitat prediction for 30,000 species”. In: *CLEF working notes* (2021).

- [12] Jian Zhang and Sen Li. “A Review of Machine Learning Based Species’ Distribution Modelling”. In: *2017 International Conference on Industrial Informatics - Computing Technology, Intelligent Technology, Industrial Information Integration (ICIICII)*. 2017, pp. 199–206. DOI: 10.1109/ICIICII.2017.76.